

# Introduction to Callbacks

Callbacks are a useful piece of functionality in Tensorflow that lets you have control over the training process, there's two main flavor of callback. There's the built in callbacks that are pre built in functions that allow you to do things like saving checkpoints early, stopping on this custom callbacks where you can override the callback class to do whatever you want.

In this section, I'm going to look at the built in callbacks, and then later you could learn how to do the custom ones.

## Callbacks

- Provides some functionality at various stages of training
- Subclasses `tf.keras.callbacks.Callback`
- Useful in understanding a model's state during training
  - internal states
  - statistics e.g., losses and metrics

So in summary, callbacks are designed to give you some type of functionality, while you're training every epoch, you can effectively have code that executes to perform a task.

What that task is is up to you, there's a `tf.keras.callbacks.Callback` class that you'll subclass, so the pattern you've been looking at in this course for subclassing existing objects will also work for this. They're particularly useful in helping you understand the model state during training, saving you valuable time is your optimizing your model.

## Training specific methods

```
class Callback(object):  
    def __init__(self):  
        self.validation_data = None  
        self.model = None
```

Here's the anatomy of a callback class, as with any class in python, you define it to extend an existing class and you have local variables initialized in the init function.

For callbacks, you have the epoch begin function that you can override, which, as its name suggests, gets called at the beginning of every epoch.

```
def on_epoch_begin(self, epoch, logs=None):  
    """Called at the beginning of an epoch during training."""
```

```
def on_epoch_end(self, epoch, logs=None):  
    """Called at the end of an epoch during training."""
```

## Training specific methods

```
class Callback(object):
    def __init__(self):
        self.validation_data = None
        self.model = None

    def on_epoch_begin(self, epoch, logs=None):
        """Called at the beginning of an epoch during training."""

    def on_epoch_end(self, epoch, logs=None):
        """Called at the end of an epoch during training."""
```

## Common methods for training/testing/predicting

```
class Callback(object):
    ...
    def on_(train|test|predict)_begin(self, logs=None):
        """Called at the begin of fit/evaluate/predict."""

    def on_(train|test|predict)_end(self, logs=None):
        """Called at the end of fit/evaluate/predict."""

    def on_(train|test|predict)_batch_begin(self, batch, logs=None):
        """Called right before processing a batch during training/testing/predicting."""

    def on_(train|test|predict)_batch_end(self, batch, logs=None):
        """Called at the end of training/testing/predicting a batch."""
```

So, for example, when you run a prediction, you can have a callback that happens at the beginning or the end by calling the `on_predict` begin method or the `on_predict` end method, respectively, you could do similar for on training and on testing.

## Common methods for training/testing/predicting

```
class Callback(object):
    ...
    def on_(train|test|predict)_begin(self, logs=None):
        """Called at the begin of fit/evaluate/predict."""

    def on_(train|test|predict)_end(self, logs=None):
        """Called at the end of fit/evaluate/predict."""

    def on_(train|test|predict)_batch_begin(self, batch, logs=None):
        """Called right before processing a batch during training/testing/predicting."""

    def on_(train|test|predict)_batch_end(self, batch, logs=None):
        """Called at the end of training/testing/predicting a batch."""
```

Also, when running batches, you can have `on_predict` batch begin or `on_predict` batch end so that you can execute code batch by batch while predicting, and of course, you could do similar for training and testing.

## Where can you use them?

Model methods that take callbacks

- `fit(..., callbacks=[...])`
- `fit_generator(..., callbacks=[...])`
- `evaluate(..., callbacks=[...])`
- `evaluate_generator(..., callbacks=[...])`
- `predict(..., callbacks=[...])`
- `predict_generator(..., callbacks=[...])`

This, of course, leads to the question, where would you use callbacks?

Well, the model methods that involve training, evaluation of prediction used them, you simply specify them, using the `callbacks` equal parameter.

## TensorBoard Callback

- Visualize machine learning experiments
- Track metrics (e.g., loss, accuracy)
- View the model graph

So now let's take a look at some of the built in callbacks, we'll start with `TensorBoard`, which, if you aren't familiar with it, provides a suite of visualization tools for tensorflow.

Unless you visualize your experiments and track metrics like loss and accuracy, Acela's viewing the model graph you can learn more about [attentive load torque slash centerboard](#)

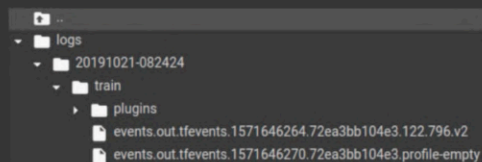
```
TensorBoard(log_dir='./logs', update_freq='epoch', **kwargs)
```

<https://www.tensorflow.org/tensorboard>

TensorBoard

Define the callback and start training

```
log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
model.fit(train_batches, epochs=10, callbacks=[tensorboard])
```




To use the `TensorBoard` callback is super simple. You simply define it and then start training, it's defined by creating an instance of the `TensorBoard` callback and specifying the desired log directory.

## TensorBoard

## Define the callback and start training

```
log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
model.fit(train_batches, epochs=10, callbacks=[tensorboard])
```



Tensorflow then saves the details to that directory

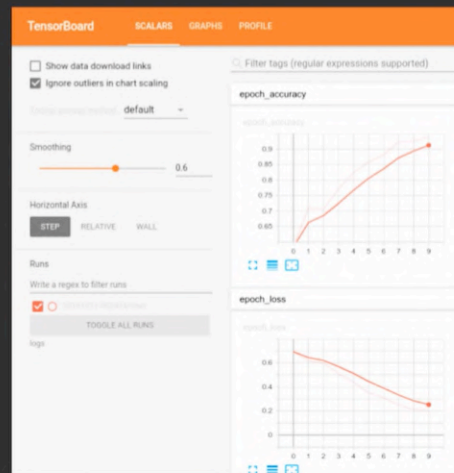
## TensorBoard

## TensorBoard in Colab

Then, when Tensorboard is pointed at the logs directory, it does its thing such as plotting accuracy and loss. It can even be used by colab by loading it as an extension, as you can see here.

```
# Load the extension
%load_ext tensorboard

# Run TensorBoard
%tensorboard --logdir logs
```



# Model Checkpoints

Next, take a look at the model checkpoints where the models details can be saved out, epoch by epoch for later inspection, or we can monitor progress through them.

# ModelCheckpoint

The model checkpoint class saves the model details for you with a lot of parameters that you can use to fine tune it.

- Saves the model every so often
- Choose to save only the best checkpoints / weights

```
ModelCheckpoint(filepath, monitor='val_loss', mode='auto',
                save_best_only=False, save_weights_only=False,
                verbose=0, save_freq=1, **kwargs)
```

## ModelCheckpoint

## Saving model checkpoints

```
model.fit(train_batches, epochs=5, validation_data=validation_batches,
          callbacks=[ModelCheckpoint('model.h5', verbose=1)])
```

Using the callbacks parameter, I specified that I want a model checkpoint with the model file being called modeled on h5.

```
Epoch 1/5
Epoch 00001: saving model to model.h5
33/33 - 7s - loss: 0.6879 - accuracy: 0.6702 - val_loss: 0.0000e+00 - val_accuracy: 0.0000e+00
Epoch 2/5
Epoch 00002: saving model to model.h5
33/33 - 6s - loss: 0.6721 - accuracy: 0.8447 - val_loss: 0.6608 - val_accuracy: 0.8667
Epoch 3/5
Epoch 00003: saving model to model.h5
33/33 - 6s - loss: 0.6435 - accuracy: 0.8840 - val_loss: 0.6217 - val_accuracy: 0.9417
Epoch 4/5
Epoch 00004: saving model to model.h5
33/33 - 6s - loss: 0.5920 - accuracy: 0.8849 - val_loss: 0.5591 - val_accuracy: 0.8667
Epoch 5/5
Epoch 00005: saving model to model.h5
33/33 - 6s - loss: 0.5047 - accuracy: 0.9089 - val_loss: 0.4485 - val_accuracy: 0.8583
<tensorflow.python.keras.callbacks.History at 0x7f09ccef97f0>
```

## ModelCheckpoint

## Saving model checkpoints

```
model.fit(train_batches, epochs=5, validation_data=validation_batches,
          callbacks=[ModelCheckpoint('model.h5', verbose=1)])
```

Then, during the training process, you can see that the model is getting saved out epoch by epoch.

```
Epoch 1/5
Epoch 00001: saving model to model.h5
33/33 - 7s - loss: 0.6879 - accuracy: 0.6702 - val_loss: 0.0000e+00 - val_accuracy: 0.0000e+00
Epoch 2/5
Epoch 00002: saving model to model.h5
33/33 - 6s - loss: 0.6721 - accuracy: 0.8447 - val_loss: 0.6608 - val_accuracy: 0.8667
Epoch 3/5
Epoch 00003: saving model to model.h5
33/33 - 6s - loss: 0.6435 - accuracy: 0.8840 - val_loss: 0.6217 - val_accuracy: 0.9417
Epoch 4/5
Epoch 00004: saving model to model.h5
33/33 - 6s - loss: 0.5920 - accuracy: 0.8849 - val_loss: 0.5591 - val_accuracy: 0.8667
Epoch 5/5
Epoch 00005: saving model to model.h5
33/33 - 6s - loss: 0.5047 - accuracy: 0.9089 - val_loss: 0.4485 - val_accuracy: 0.8583
<tensorflow.python.keras.callbacks.History at 0x7f09ccef97f0>
```



## ModelCheckpoint

## Save only the weights

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,
        callbacks=[ModelCheckpoint('model.h5', save_weights_only=True, verbose=1)])
```

If I don't want the entire model structure and only the weights, I could do so by specifying the save\_weights\_only parameter to be true.

Epoch 1/2

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6493 - accuracy: 0.6184 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/2

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.5684 - accuracy: 0.7507 - val\_loss: 0.5183 - val\_accuracy: 0.7083

<tensorflow.python.keras.callbacks.History at 0x7f09cb5547f0>

## ModelCheckpoint

## Save only the best checkpoints

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,
        callbacks=[ModelCheckpoint('model.h5', monitor='val_loss',
                                   save_best_only=True, verbose=1)])
```

Epoch 1/5

Epoch 00001: val loss improved from inf to 0.65278, saving model to model.h5

33/33 - 7s - loss: 0.6753 - accuracy: 0.5772 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: val loss improved from 0.65278 to 0.62279, saving model to model.h5

33/33 - 6s - loss: 0.6219 - accuracy: 0.7584 - val\_loss: 0.6228 - val\_accuracy: 0.5417

Epoch 3/5

Epoch 00003: val loss improved from 0.62279 to 0.47633, saving model to model.h5

33/33 - 6s - loss: 0.5448 - accuracy: 0.7977 - val\_loss: 0.4763 - val\_accuracy: 0.8750

Epoch 4/5

Epoch 00004: val loss improved from 0.47633 to 0.44497, saving model to model.h5

33/33 - 6s - loss: 0.4673 - accuracy: 0.8054 - val\_loss: 0.4450 - val\_accuracy: 0.8000

Epoch 5/5

Epoch 00005: val loss improved from 0.44497 to 0.30997, saving model to model.h5

33/33 - 6s - loss: 0.4030 - accuracy: 0.8677 - val\_loss: 0.3100 - val\_accuracy: 0.9000

<tensorflow.python.keras.callbacks.History at 0x7f09cc9b7128>

Or if I only want to save when I reach optimal values, I could do so by specifying save\_best\_only to be true, then whatever value I specify in the monitor parameter will be saved whenever it's optimized. As you can see here in the first epoch, the value started as infinite, and it ended at 0.65278 so get safe

## ModelCheckpoint

## Save only the best checkpoints

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,
        callbacks=[ModelCheckpoint('model.h5', monitor='val_loss',
                                   save_best_only=True, verbose=1)])
```

Epoch 1/5

Epoch 00001: val loss improved from inf to 0.65278, saving model to model.h5

33/33 - 7s - loss: 0.6753 - accuracy: 0.5772 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: val loss improved from 0.65278 to 0.62279, saving model to model.h5

33/33 - 6s - loss: 0.6219 - accuracy: 0.7584 - val\_loss: 0.6228 - val\_accuracy: 0.5417

Epoch 3/5

Epoch 00003: val loss improved from 0.62279 to 0.47633, saving model to model.h5

33/33 - 6s - loss: 0.5448 - accuracy: 0.7977 - val\_loss: 0.4763 - val\_accuracy: 0.8750

Epoch 4/5

Epoch 00004: val loss improved from 0.47633 to 0.44497, saving model to model.h5

33/33 - 6s - loss: 0.4673 - accuracy: 0.8054 - val\_loss: 0.4450 - val\_accuracy: 0.8000

Epoch 5/5

Epoch 00005: val loss improved from 0.44497 to 0.30997, saving model to model.h5

33/33 - 6s - loss: 0.4030 - accuracy: 0.8677 - val\_loss: 0.3100 - val\_accuracy: 0.9000

<tensorflow.python.keras.callbacks.History at 0x7f09cc9b7128>

Then in the second epoch, it improved, so it got saved and so on. If at some point you're val loss starts to increase, the model checkpoints, of course, would not be saved.

ModelCheckpoint

Choose your model format (SavedModel / H5)

```
model.fit(..., callbacks=[ModelCheckpoint('saved_model', ...)])
```

```
Epoch 1/2
Epoch 00001: saving model to model.h5
33/33 - 7s - loss: 0.6714 - accuracy: 0.5695 - val_loss: 0.0000e+00 - val_accuracy: 0.0000e+00
Epoch 2/2
Epoch 00002: saving model to model.h5
33/33 - 6s - loss: 0.6238 - accuracy: 0.6366 - val_loss: 0.6459 - val_accuracy: 0.5417
```

```
model.fit(..., callbacks=[ModelCheckpoint('model.h5', ...)])
```

```
Epoch 1/2
Epoch 00001: saving model to model.h5
33/33 - 7s - loss: 0.6714 - accuracy: 0.5695 - val_loss: 0.0000e+00 - val_accuracy: 0.0000e+00
Epoch 2/2
Epoch 00002: saving model to model.h5
33/33 - 6s - loss: 0.6238 - accuracy: 0.6366 - val_loss: 0.6459 - val_accuracy: 0.5417
```

In these examples, I've been showing the native keras h5 format, but of course, you can also use a save\_model, which is the standard tensorflow format.

```

content
├── saved_model
├── assets
├── variables
│   ├── variables.data-00000-of-00002
│   ├── variables.data-00001-of-00002
│   ├── variables.index
│   └── saved_model.pb
└── model.h5

```

ModelCheckpoint

Track epoch #, losses, metrics

```
model.fit(..., callbacks=[ModelCheckpoint('weights.{epoch:02d}-{val_loss:.2f}.h5', verbose=1)])
```

```
Epoch 1/5
Epoch 00001: saving model to weights.01-0.63.h5
33/33 - 6s - loss: 0.6709 - accuracy: 0.6098 - val_loss: 0.0000e+00 - val_accuracy: 0.0000e+00
Epoch 2/5
Epoch 00002: saving model to weights.02-0.60.h5
33/33 - 6s - loss: 0.6088 - accuracy: 0.7124 - val_loss: 0.6046 - val_accuracy: 0.5917
Epoch 3/5
Epoch 00003: saving model to weights.03-0.46.h5
33/33 - 6s - loss: 0.5354 - accuracy: 0.7613 - val_loss: 0.4602 - val_accuracy: 0.8500
Epoch 4/5
Epoch 00004: saving model to weights.04-0.38.h5
33/33 - 6s - loss: 0.4769 - accuracy: 0.7891 - val_loss: 0.3848 - val_accuracy: 0.9250
Epoch 5/5
Epoch 00005: saving model to weights.05-0.33.h5
33/33 - 6s - loss: 0.3961 - accuracy: 0.8600 - val_loss: 0.3263 - val_accuracy: 0.8667
```

```

content
├── weights.01-0.63.h5
├── weights.02-0.60.h5
├── weights.03-0.46.h5
├── weights.04-0.38.h5
└── weights.05-0.33.h5

```

As the name of the file is just specified using text, you can actually form at the values within the name, so you could have separate weights saved out pair epoch in separate h5 files. Simply by using the epoch value or other metrics such as the validation lost value you conform at the file name so you can see here the last two digits of the epoch are used, so the file is wait 01 wait 02 and so on.

## EarlyStopping

- Helps you keep track of a certain metric/loss and change training behavior accordingly
- Stops training when there's no improvement observed

```
EarlyStopping(monitor='val_loss', verbose=0,
              min_delta=0, patience=0, mode='auto',
              baseline=None, restore_best_weights=False,
              **kwargs)
```

## EarlyStopping

## Prevent your model from overfitting

Set patients to three, the idea here is that once we hit the best value will log that and we'll wait for this number of epochs ie three to see if the values improve.

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500
Epoch 12/50
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000
Epoch 13/50
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500
Epoch 14/50
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583
Epoch 15/50
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583
Epoch 16/50
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583
Epoch 17/50
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167
Epoch 18/50
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333
Epoch 00018: early stopping
```

## EarlyStopping

## Restoring best weights

If you don't want to lose the weight values from the best epoch you can set, restore best weights to be true. So in our case, even though we stopped at 18 will have the weight restored to where they were at 15.

```
model.fit(...,
          callbacks=[EarlyStopping(patience=3, restore_best_weights=True,
                                   monitor='val_loss', verbose=1)])
```

```
Epoch 11/50
33/33 - 6s - loss: 0.1380 - accuracy: 0.9616 - val_loss: 0.0968 - val_accuracy: 0.9750
Epoch 12/50
33/33 - 6s - loss: 0.1202 - accuracy: 0.9655 - val_loss: 0.0741 - val_accuracy: 0.9917
Epoch 13/50
33/33 - 6s - loss: 0.1716 - accuracy: 0.9434 - val_loss: 0.1083 - val_accuracy: 0.9750
Epoch 14/50
33/33 - 6s - loss: 0.1331 - accuracy: 0.9626 - val_loss: 0.0861 - val_accuracy: 0.9667
Epoch 15/50
33/33 - 6s - loss: 0.1393 - accuracy: 0.9578 - val_loss: 0.0771 - val_accuracy: 0.9750
Epoch 00015: early stopping
```

## EarlyStopping

## More customization

```
model.fit(...,
          callbacks=[EarlyStopping(
              patience=3,
              min_delta=0.05,
              baseline=0.8,
              mode='min',
              monitor='val_loss',
              verbose=1
          )])
```

There's other parameters you can play with, but the mode one is crucial to ensure that you're following your monitor values correctly for loss that you want to minimize. So you would then set the mode to min, for others, they might require you to maximize the value so you could change the mode with this property.



```
model.fit(..., callbacks=[CSVLogger('training.csv')])
```

epoch	accuracy	loss	val_accuracy	val_loss
0	0.574305	0.682536	0.775000	0.655427
1	0.760307	0.633610	0.675000	0.595201
2	0.758389	0.573186	0.850000	0.503174
3	0.835091	0.472031	0.808333	0.416691
4	0.854267	0.419491	0.916667	0.309128

Another super useful callback is the CSVLogger which, as its name suggests, will log your training results out to a CSV file. So, for example, when using it like this, you'll have a file containing the epoch number, accuracy, loss, validation accuracy and validation loss stored for you.

## Build a simple model

```
model = tf.keras.Sequential()

model.add(tf.keras.layers.Dense(units=1,
                                activation='linear',
                                input_dim=(784,)))

model.compile(optimizer=tf.keras.optimizers.RMSprop(lr=0.1),
              loss='mean_squared_error', metrics=['mae'])
```

## How a custom callback looks

```
import datetime

class MyCustomCallback(tf.keras.callbacks.Callback):
    def on_train_batch_begin(self, batch, logs=None):
        print('Training: batch {} begins at {}'.format(batch, datetime.datetime.now().time()))

    def on_train_batch_end(self, batch, logs=None):
        print('Training: batch {} ends at {}'.format(batch, datetime.datetime.now().time()))
```

Now let's define a custom callback, which we'll name my custom callback. We'll do this by subclassing the `tf.keras.callbacks.CallbackClass`.

Our custom callback will print the timestamp each time when a training batch starts or ends. To achieve this we'll override the `on_train_batch_begin` and `on_train_batch_end` methods to include our business logic, which just displays the timestamps.

To use the custom callback that we've just defined we'll instantiate an instance of this class. Here we save the instance of the custom callback in a variable named `my_custom_callback` and as usual, we'll train our model using `model.fit`.

```
my_custom_callback = MyCustomCallback()
```

```
model.fit(x_train, y_train, batch_size=64, epochs=1, verbose=0,
         callbacks=[my_custom_callback])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):
```

```
    def __init__(self, threshold):
```

```
        super(DetectOverfittingCallback, self).__init__()
```

```
        self.threshold = threshold
```

```
    def on_epoch_end(self, epoch, logs=None):
```

```
        ratio = logs["val_loss"] / logs["loss"]
```

```
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))
```

```
        if ratio > threshold:
```

```
            print("Stopping training...")
```

```
            self.model.stop_training = True
```

```
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

To use the custom callback that we've just defined we'll instantiate an instance of this class. Here we save the instance of the custom callback in a variable named `my_custom_callback` and as usual, we'll train our model using `model.fit`.

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):
```

```
    def __init__(self, threshold):
```

```
        super(DetectOverfittingCallback, self).__init__()
```

```
        self.threshold = threshold
```

```
    def on_epoch_end(self, epoch, logs=None):
```

```
        ratio = logs["val_loss"] / logs["loss"]
```

```
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))
```

```
        if ratio > threshold:
```

```
            print("Stopping training...")
```

```
            self.model.stop_training = True
```

```
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

Let's explore a call back where we measure the ratio between our validation loss and our training loss.

If the ratio gets too high, we could have an over-fitting scenario because the validation loss may no longer be decreasing while the training loss continues to decrease, making the ratio of validation loss divided by training loss higher.

We should in this case, stop training to avoid overfitting. We'll do this by defining a class called Detect Overfitting Callback and this subclass is the Keras callback base class.

We'll need a class level variable to hold the threshold so that we can override the init function to take the threshold as a parameter and then store it in `self.threshold`.

```

class DetectOverfittingCallback(tf.keras.callbacks.Callback):
    def __init__(self, threshold):
        super(DetectOverfittingCallback, self).__init__()
        self.threshold = threshold

    def on_epoch_end(self, epoch, logs=None):
        ratio = logs["val_loss"] / logs["loss"]
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))

        if ratio > threshold:
            print("Stopping training...")
            self.model.stop_training = True

model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])

```

We can then implement the on epoch end class method overriding the base class. In this method we'll compute the ratio at the end of every epoch.

```

class DetectOverfittingCallback(tf.keras.callbacks.Callback):
    def __init__(self, threshold):
        super(DetectOverfittingCallback, self).__init__()
        self.threshold = threshold

    def on_epoch_end(self, epoch, logs=None):
        ratio = logs["val_loss"] / logs["loss"]
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))

        if ratio > threshold:
            print("Stopping training...")
            self.model.stop_training = True

model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])

```

If that ratio was higher than our threshold value, we can stop training.

```

class DetectOverfittingCallback(tf.keras.callbacks.Callback):
    def __init__(self, threshold):
        super(DetectOverfittingCallback, self).__init__()
        self.threshold = threshold

    def on_epoch_end(self, epoch, logs=None):
        ratio = logs["val_loss"] / logs["loss"]
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))

        if ratio > threshold:
            print("Stopping training...")
            self.model.stop_training = True

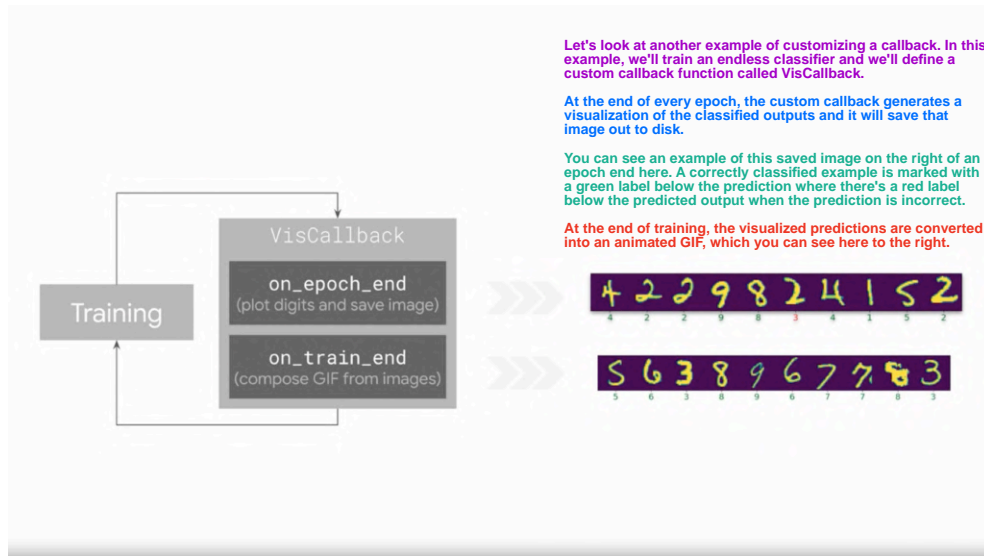
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])

```

To use this, we'll just then specify detect overfitting callback when performing the model.fit and pass in our desired threshold.

Notice that we're creating an instance of the class by invoking the detect overfitting callback class constructor and placing it directly into the list that gets passed to the callbacks parameter.

You could have also stored an instance of the callback class in a variable and passed in that variable, as you saw in the earlier example with my custom callback.



```
class VisCallback(tf.keras.callbacks.Callback):
    def __init__(self, inputs, ground_truth, display_freq=10,
                  n_samples=10):

        self.inputs = inputs
        self.ground_truth = ground_truth
        self.images = []
        self.display_freq = display_freq
        self.n_samples = n_samples
```

First we'll define a custom callback called VisCallback subclassing `tf.keras.callbacks.Callback` and we'll have a set of class variables.

`self.inputs` holds the inputs to run in the model.

`self.ground_truth` holds the ground truth or the true labels for those input images. That allows us to compare each prediction with the actual label.

```
class VisCallback(tf.keras.callbacks.Callback):
    def __init__(self, inputs, ground_truth, display_freq=10,
                  n_samples=10):

        self.inputs = inputs
        self.ground_truth = ground_truth
        self.images = []
        self.display_freq = display_freq
        self.n_samples = n_samples
```

`Self.images` will hold the visualize comparison of the prediction against the ground truth.

```

class VisCallback(tf.keras.callbacks.Callback):
    def __init__(self, inputs, ground_truth, display_freq=10,
                  n_samples=10):

        self.inputs = inputs
        self.ground_truth = ground_truth
        self.images = []
        self.display_freq = display_freq
        self.n_samples = n_samples

```

Self.display frequency allows you to choose how frequently you want to display these plots. For instance, you might decide to display a visual at every 10 epochs instead of every epoch.

```

class VisCallback(tf.keras.callbacks.Callback):
    def __init__(self, inputs, ground_truth, display_freq=10,
                  n_samples=10):

        self.inputs = inputs
        self.ground_truth = ground_truth
        self.images = []
        self.display_freq = display_freq
        self.n_samples = n_samples

```

Self.n\_samples then determines how many samples to be plotted each time a visualization is generated.

```

class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_epoch_end(self, epoch, logs=None):
        # Randomly sample data
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]
        predictions = np.argmax(self.model.predict(X_test), axis=1)

```

At the end of every epoch we'll randomly sample data from the list of input images and then classify them.

Accordingly, we'll mark the label predictions green or red based on whether they're classified correctly or not.

We'll start by using np.random.choice passing it the number of inputs and the size of samples that we want. This will give us a list of random indexes back and we can use these to pick the images from our training data completely at random.



```

class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_epoch_end(self, epoch, logs=None):
        # Randomly sample data
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]
        predictions = np.argmax(self.model.predict(X_test), axis=1)

```

We can then get our x values from the test, from the inputs and using those randomly generated indexes, the corresponding y values from the ground truth.

```

class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_epoch_end(self, epoch, logs=None):
        # Randomly sample data
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]
        predictions = np.argmax(self.model.predict(X_test), axis=1)

```

To get our predictions, you can use `model.predict` and pass it the X test values. You'll use `argmax` to get the most likely value.

Recall that `argmax` will take a list, find the maximum value on that list, and then return the index position of that maximum value.

```

class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_epoch_end(self, epoch, logs=None):
        # Randomly sample data
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]
        predictions = np.argmax(self.model.predict(X_test), axis=1)

        # Plot the digits
        display_digits(X_test, predictions, y_test, epoch, n=self.display_freq)
    ...

```

Now that we have our predictions, we can pass our data to a display function to draw them in our plotter.

```

class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_epoch_end(self, epoch, logs=None):
        ...
        # Save the figure
        buf = io.BytesIO()
        plt.savefig(buf, format='png')
        buf.seek(0)
        image = Image.open(buf)
        self.images.append(np.array(image))

        # Display the digits every now and then
        if epoch % self.display_freq == 0:
            plt.show()

```

Then when we save those plotted images in a list by reading the results of the plot into a buffer, which can then be appended to the list called self.images.

```

class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_epoch_end(self, epoch, logs=None):
        ...
        # Save the figure
        buf = io.BytesIO()
        plt.savefig(buf, format='png')
        buf.seek(0)
        image = Image.open(buf)
        self.images.append(np.array(image))

        # Display the digits every now and then
        if epoch % self.display_freq == 0:
            plt.show()

```

If we want to show the plots occasionally, we can do so with code like this, where for example, if the frequency is 10, it will render every tenth epoch.

Here's how you can create an animated GIF that's composed of several saved images. The imageio library for Python contains a mimsave method that lets you specify an array of images and it will write out an animated GIF of them. As you've kept a list of images in self.images, this is now really easy to do.

```

import imageio

class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_train_end(self, logs=None):
        imageio.mimsave('animation.gif', self.images, fps=1)

# Train the model
model.fit(..., callbacks=[VisCallback(x_test, y_test)])

```

```
import imageio

class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_train_end(self, logs=None):
        imageio.mimsave('animation.gif', self.images, fps=1)

# Train the model
model.fit(..., callbacks=[VisCallback(x_test, y_test)])
```

To get this visualization, you specify this custom class in the callbacks when you call your model.fit.

You'll pass in the x test and y test so that the visualizations are created on the test set and not on the training set.

### ***TensorBoard: TensorFlow's visualization toolkit***

- <https://www.tensorflow.org/tensorboard>