

Custom Training Basics

In this ungraded lab you'll gain a basic understanding of building custom training loops.

- It takes you through the underlying logic of fitting any model to a set of inputs and outputs.
- You will be training your model on the linear equation for a straight line, $wx + b$.
- You will implement basic linear regression from scratch using gradient tape.
- You will try to minimize the loss incurred by the model using linear regression.

Imports

In [1]:

```
from __future__ import absolute_import, division, print_function, unicode_literals

try:
    # %tensorflow_version only exists in Colab.
    %tensorflow_version 2.x
except Exception:
    pass

import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
```

Define Model

You define your model as a class.

- x is your input tensor.
- The model should output values of $wx+b$.
- You'll start off by initializing w and b to random values.
- During the training process, values of w and b get updated in accordance with linear regression so as to minimize the loss incurred by the model.
- Once you arrive at optimal values for w and b , the model would have been trained to correctly predict the values of $wx+b$.

Hence,

- w and b are trainable weights of the model.
- x is the input
- $y = wx + b$ is the output

In [2]:

```
class Model(object):
    def __init__(self):
        # Initialize the weights to `2.0` and the bias to `1.0`
        # In practice, these should be initialized to random values (for example, with
        `tf.random.normal`)
        self.w = tf.Variable(2.0)
        self.b = tf.Variable(1.0)

    def __call__(self, x):
        return self.w * x + self.b

model = Model()
```

Define a loss function

A loss function measures how well the output of a model for a given input matches the target output.

- The goal is to minimize this difference during training.

- Let's use the standard L2 loss, also known as the least square errors

$$Loss = \sum_i \left(y_{pred}^i - y_{target}^i \right)^2$$

In [3]:

```
def loss(predicted_y, target_y):
    return tf.reduce_mean(tf.square(predicted_y - target_y))
```

Obtain training data

First, synthesize the training data using the "true" w and "true" b .

$$y = w_{true} \times x + b_{true}$$

In [4]:

```
TRUE_w = 3.0
TRUE_b = 2.0
NUM_EXAMPLES = 1000

xs = tf.random.normal(shape=[NUM_EXAMPLES])

ys = (TRUE_w * xs) + TRUE_b
```


Before training the model, visualize the loss value by plotting the model's predictions in red crosses and the training data in blue dots:

In [5]:

```
def plot_data(inputs, outputs, predicted_outputs):
    real = plt.scatter(inputs, outputs, c='b', marker='.')
    predicted = plt.scatter(inputs, predicted_outputs, c='r', marker='+')
    plt.legend((real, predicted), ('Real Data', 'Predicted Data'))
    plt.show()
```

In [6]:

```
plot_data(xs, ys, model(xs))
print('Current loss: %1.6f' % loss(model(xs), ys).numpy())
```



Current loss: 1.819981

Define a training loop

With the network and training data, train the model using [gradient descent](#)

- Gradient descent updates the trainable weights w and b to reduce the loss.

There are many variants of the gradient descent scheme that are captured in `tf.train.Optimizer`—our recommended

implementation. In the spirit of building from first principles, here you will implement the basic math yourself.

- You'll use `tf.GradientTape` for automatic differentiation
- Use `tf.assign_sub` for decrementing a value. Note that `assign_sub` combines `tf.assign` and `tf.sub`

In [7]:

```
def train(model, inputs, outputs, learning_rate):
    with tf.GradientTape() as t:
        current_loss = loss(model(inputs), outputs)
        dw, db = t.gradient(current_loss, [model.w, model.b])
        model.w.assign_sub(learning_rate * dw)
        model.b.assign_sub(learning_rate * db)

    return current_loss
```

Finally, you can iteratively run through the training data and see how `w` and `b` evolve.

In [8]:

```
model = Model()


# Collect the history of W-values and b-values to plot later
list_w, list_b = [], []
epochs = range(15)
losses = []
for epoch in epochs:
    list_w.append(model.w.numpy())
    list_b.append(model.b.numpy())
    current_loss = train(model, xs, ys, learning_rate=0.1)
    losses.append(current_loss)
    print('Epoch %2d: w=%1.2f b=%1.2f, loss=%2.5f' %
          (epoch, list_w[-1], list_b[-1], current_loss))
```


```
Epoch 0: w=2.00 b=1.00, loss=1.81998
Epoch 1: w=2.17 b=1.19, loss=1.21552
Epoch 2: w=2.31 b=1.35, loss=0.81262
Epoch 3: w=2.43 b=1.48, loss=0.54379
Epoch 4: w=2.52 b=1.58, loss=0.36423
Epoch 5: w=2.61 b=1.66, loss=0.24418
Epoch 6: w=2.67 b=1.73, loss=0.16385
Epoch 7: w=2.73 b=1.78, loss=0.11003
Epoch 8: w=2.77 b=1.82, loss=0.07396
Epoch 9: w=2.81 b=1.86, loss=0.04975
Epoch 10: w=2.85 b=1.88, loss=0.03349
Epoch 11: w=2.87 b=1.91, loss=0.02256
Epoch 12: w=2.89 b=1.92, loss=0.01521
Epoch 13: w=2.91 b=1.94, loss=0.01026
Epoch 14: w=2.93 b=1.95, loss=0.00693
```

In addition to the values for losses, you also plot the progression of trainable variables over epochs.

In [9]:

```
plt.plot(epochs, list_w, 'r',
         epochs, list_b, 'b')
plt.plot([TRUE_w] * len(epochs), 'r--',
         [TRUE_b] * len(epochs), 'b--')
plt.legend(['w', 'b', 'True w', 'True b'])
plt.show()
```





Plots for Evaluation


Now you can plot the actual outputs in red and the model's predictions in blue on a set of random test examples.

You can see that the model is able to make predictions on the test set fairly accurately.

In [10]:

```
test_inputs = tf.random.normal(shape=[NUM_EXAMPLES])
test_outputs = test_inputs * TRUE_w + TRUE_b

predicted_test_outputs = model(test_inputs)
plot_data(test_inputs, test_outputs, predicted_test_outputs)
```




Visualize the cost function against the values of each of the trainable weights the model approximated to over time.

In [11]:

```
def plot_loss_for_weights(weights_list, losses):
    for idx, weights in enumerate(weights_list):
        plt.subplot(120 + idx + 1)
        plt.plot(weights['values'], losses, 'r')
        plt.plot(weights['values'], losses, 'bo')
        plt.xlabel(weights['name'])
        plt.ylabel('Loss')

weights_list = [{ 'name' : "w",
                  'values' : list_w },
                { 'name' : "b",
                  'values' : list_b }]

plot_loss_for_weights(weights_list, losses)
```



2.00	2.25	2.50	2.75	1.00	1.25	1.50	1.75
		w				b	