

Horse or Human? In-graph training loop Assignment

This assignment lets you practice how to train a Keras model on the [horses_or_humans](#) dataset with the entire training process performed in graph mode. These steps include:

- loading batches
- calculating gradients
- updating parameters
- calculating validation accuracy
- repeating the loop until convergence

Setup

Import TensorFlow 2.0:

In [1]:

```
from __future__ import absolute_import, division, print_function, unicode_literals
import numpy as np
```

In [2]:

```
import tensorflow as tf
import tensorflow_datasets as tfds
import tensorflow_hub as hub
import matplotlib.pyplot as plt
```

Prepare the dataset

Load the horses to human dataset, splitting 80% for the training set and 20% for the test set.

In [3]:

```
splits, info = tfds.load('horses_or_humans', as_supervised=True, with_info=True, split=['train[:80%]', 'train[80%:]', 'test'], data_dir='./data')

(train_examples, validation_examples, test_examples) = splits

num_examples = info.splits['train'].num_examples
num_classes = info.features['label'].num_classes
```

In [4]:

```
BATCH_SIZE = 32
IMAGE_SIZE = 224
```

Pre-process an image (please complete this section)

You'll define a mapping function that resizes the image to a height of 224 by 224, and normalizes the pixels to the range of 0 to 1. Note that pixels range from 0 to 255.

- You'll use the following function: [tf.image.resize](#) and pass in the (height,width) as a tuple (or list).
- To normalize, divide by a floating value so that the pixel range changes from [0,255] to [0,1].

In [5]:

```
# Create a autograph pre-processing function to resize and normalize an image
### START CODE HERE ###
@tf.function
def map_fn(img, label):
    image_height = 224
    image_width = 224
```

```

    image_width = 224
    ### START CODE HERE ###
    # resize the image
    img = tf.image.resize(img, (image_height, image_width))
    # normalize the image
    img /= 255
    ### END CODE HERE
    return img, label

```

In [6]:

```

## TEST CODE:

test_image, test_label = list(train_examples)[0]

test_result = map_fn(test_image, test_label)

print(test_result[0].shape)
print(test_result[1].shape)

del test_image, test_label, test_result

```

```

(224, 224, 3)
()

```

Expected Output:

```

(224, 224, 3)
()

```

Apply pre-processing to the datasets (please complete this section)

Apply the following steps to the training_examples:

- Apply the `map_fn` to the training_examples
- Shuffle the training data using `.shuffle(buffer_size=)` and set the buffer size to the number of examples.
- Group these into batches using `.batch()` and set the batch size given by the parameter.

Hint: You can look at how validation_examples and test_examples are pre-processed to get a sense of how to chain together multiple function calls.

In [7]:

```

# Prepare train dataset by using preprocessing with map_fn, shuffling and batching
def prepare_dataset(train_examples, validation_examples, test_examples, num_examples, map_fn, batch_size):
    ### START CODE HERE ###
    train_ds = train_examples.map(map_fn).shuffle(buffer_size=num_examples).batch(batch_size)
    ### END CODE HERE ###
    valid_ds = validation_examples.map(map_fn).batch(batch_size)
    test_ds = test_examples.map(map_fn).batch(batch_size)

    return train_ds, valid_ds, test_ds

```

In [8]:

```

train_ds, valid_ds, test_ds = prepare_dataset(train_examples, validation_examples, test_examples, num_examples, map_fn, BATCH_SIZE)

```

In [9]:

```

## TEST CODE:

test_train_ds = list(train_ds)
print(len(test_train_ds))
print(test_train_ds[0][0].shape)

del test_train_ds

```

```
26
(32, 224, 224, 3)
```

Expected Output:

```
26
(32, 224, 224, 3)
```

Define the model

In [10]:

```
MODULE_HANDLE = 'data/resnet_50_feature_vector'
model = tf.keras.Sequential([
    hub.KerasLayer(MODULE_HANDLE, input_shape=(IMAGE_SIZE, IMAGE_SIZE, 3)),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
keras_layer (KerasLayer)	(None, 2048)	23561152
dense (Dense)	(None, 2)	4098

Total params: 23,565,250
Trainable params: 4,098
Non-trainable params: 23,561,152

Define optimizer: (please complete these sections)

Define the [Adam optimizer](#) that is in the `tf.keras.optimizers` module.

In [11]:

```
def set_adam_optimizer():
    ### START CODE HERE ###
    # Define the adam optimizer
    optimizer = tf.keras.optimizers.Adam()
    ### END CODE HERE ###
    return optimizer
```

In [12]:

```
## TEST CODE:

test_optimizer = set_adam_optimizer()

print(type(test_optimizer))

del test_optimizer
```

```
<class 'tensorflow.python.keras.optimizer_v2.adam.Adam'>
```

Expected Output:

```
<class 'tensorflow.python.keras.optimizer_v2.adam.Adam'>
```

Define the loss function (please complete this section)

Define the loss function as the [sparse categorical cross entropy](#) that's in the `tf.keras.losses` module. Use the same function for both training and validation.

In [13]:

```
def set_sparse_cat_crossentropy_loss():
    ### START CODE HERE ###
    # Define object oriented metric of Sparse categorical crossentropy for train and val loss
    train_loss = tf.keras.losses.SparseCategoricalCrossentropy()
    val_loss = tf.keras.losses.SparseCategoricalCrossentropy()
    ### END CODE HERE ###
    return train_loss, val_loss
```

In [14]:

```
## TEST CODE:

test_train_loss, test_val_loss = set_sparse_cat_crossentropy_loss()

print(type(test_train_loss))
print(type(test_val_loss))

del test_train_loss, test_val_loss
```

```
<class 'tensorflow.python.keras.losses.SparseCategoricalCrossentropy'>
<class 'tensorflow.python.keras.losses.SparseCategoricalCrossentropy'>
```

Expected Output:

```
<class 'tensorflow.python.keras.losses.SparseCategoricalCrossentropy'>
<class 'tensorflow.python.keras.losses.SparseCategoricalCrossentropy'>
```

Define the accuracy function (please complete this section)

Define the accuracy function as the [sparse categorical accuracy](#) that's contained in the `tf.keras.metrics` module. Use the same function for both training and validation.

In [15]:

```
def set_sparse_cat_crossentropy_accuracy():
    ### START CODE HERE ###
    # Define object oriented metric of Sparse categorical accuracy for train and val accuracy
    train_accuracy = tf.keras.metrics.SparseCategoricalAccuracy()
    val_accuracy = tf.keras.metrics.SparseCategoricalAccuracy()
    ### END CODE HERE ###
    return train_accuracy, val_accuracy
```

In [16]:

```
## TEST CODE:

test_train_accuracy, test_val_accuracy = set_sparse_cat_crossentropy_accuracy()

print(type(test_train_accuracy))
print(type(test_val_accuracy))

del test_train_accuracy, test_val_accuracy
```

```
<class 'tensorflow.python.keras.metrics.SparseCategoricalAccuracy'>
<class 'tensorflow.python.keras.metrics.SparseCategoricalAccuracy'>
```

Expected Output:

```
<class 'tensorflow.python.keras.metrics.SparseCategoricalAccuracy'>
<class 'tensorflow.python.keras.metrics.SparseCategoricalAccuracy'>
```

Call the three functions that you defined to set the optimizer, loss and accuracy

In [17]:

```
optimizer = set_adam_optimizer()
train_loss, val_loss = set_sparse_cat_crossentropy_loss()
train_accuracy, val_accuracy = set_sparse_cat_crossentropy_accuracy()
```

Define the training loop (please complete this section)

In the training loop:

- Get the model predictions: use the model, passing in the input `x`
- Get the training loss: Call `train_loss`, passing in the true `y` and the predicted `y`.
- Calculate the gradient of the loss with respect to the model's variables: use `tape.gradient` and pass in the loss and the model's `trainable_variables`.
- Optimize the model variables using the gradients: call `optimizer.apply_gradients` and pass in a `zip()` of the two lists: the gradients and the model's `trainable_variables`.
- Calculate accuracy: Call `train_accuracy`, passing in the true `y` and the predicted `y`.

In [18]:

```
# this code uses the GPU if available, otherwise uses a CPU
device = '/gpu:0' if tf.test.is_gpu_available() else '/cpu:0'
EPOCHS = 2

# Custom training step
def train_one_step(model, optimizer, x, y, train_loss, train_accuracy):
    """
    Trains on a batch of images for one step.

    Args:
        model (keras Model) -- image classifier
        optimizer (keras Optimizer) -- optimizer to use during training
        x (Tensor) -- training images
        y (Tensor) -- training labels
        train_loss (keras Loss) -- loss object for training
        train_accuracy (keras Metric) -- accuracy metric for training
    """
    with tf.GradientTape() as tape:
        ### START CODE HERE ###
        # Run the model on input x to get predictions
        predictions = model(x)
        # Compute the training loss using `train_loss`, passing in the true y and the predicted y
        loss = train_loss(y, predictions)

        # Using the tape and loss, compute the gradients on model variables using tape.gradient
        grads = tape.gradient(loss, model.trainable_weights)

        # Zip the gradients and model variables, and then apply the result on the optimizer
        optimizer.apply_gradients(zip(grads, model.trainable_weights))

        # Call the train accuracy object on ground truth and predictions
        train_accuracy.update_state(y, predictions)
    ### END CODE HERE
    return loss
```

WARNING:tensorflow:From <ipython-input-18-3906833a8de7>:2: is_gpu_available (from tensorflow.python.framework.test_util) is deprecated and will be removed in a future version. Instructions for updating:
Use `tf.config.list_physical_devices('GPU')` instead.

WARNING:tensorflow:From <ipython-input-18-3906833a8de7>:2: is_gpu_available (from tensorflow.python.framework.test_util) is deprecated and will be removed in a future version. Instructions for updating:
Use `tf.config.list_physical_devices('GPU')` instead.

In [19]:

```
## TEST CODE:

def base_model():
    inputs = tf.keras.layers.Input(shape=(2))
    x = tf.keras.layers.Dense(64, activation='relu')(inputs)
    outputs = tf.keras.layers.Dense(1, activation='sigmoid')(x)
    model = tf.keras.Model(inputs=inputs, outputs=outputs)
    return model

test_model = base_model()

test_optimizer = set_adam_optimizer()
test_image = tf.ones((2,2))
test_label = tf.ones((1,))
test_train_loss, _ = set_sparse_cat_crossentropy_loss()
test_train_accuracy, _ = set_sparse_cat_crossentropy_accuracy()

test_result = train_one_step(test_model, test_optimizer, test_image, test_label, test_train_loss, test_train_accuracy)
print(test_result)

def test_result, test_model, test_optimizer, test_image, test_label, test_train_loss, test_train_accuracy:

tf.Tensor(0.6931472, shape=(), dtype=float32)
```

Expected Output:

You will see a Tensor with the same shape and dtype. The value might be different.

```
tf.Tensor(0.6931472, shape=(), dtype=float32)
```

Define the 'train' function (please complete this section)

You'll first loop through the training batches to train the model. (Please complete these sections)

- The `train` function will use a for loop to iteratively call the `train_one_step` function that you just defined.
- You'll use `tf.print` to print the step number, loss, and `train_accuracy.result()` at each step. Remember to use `tf.print` when you plan to generate autograph code.

Next, you'll loop through the batches of the validation set to calculate the validation loss and validation accuracy. (This code is provided for you). At each iteration of the loop:

- Use the model to predict on `x`, where `x` is the input from the validation set.
- Use `val_loss` to calculate the validation loss between the true validation 'y' and predicted y.
- Use `val_accuracy` to calculate the accuracy of the predicted y compared to the true y.

Finally, you'll print the validation loss and accuracy using `tf.print`. (Please complete this section)

- print the final `loss`, which is the validation loss calculated by the last loop through the validation dataset.
- Also print the `val_accuracy.result()`.

HINT If you submit your assignment and see this error for your stderr output:

```
Cannot convert 1e-07 to EagerTensor of dtype int64
```

Please check your calls to `train_accuracy` and `val_accuracy` to make sure that you pass in the true and predicted values in the correct order (check the documentation to verify the order of parameters).

In [20]:

```
# Decorate this function with tf.function to enable autograph on the training loop
@tf.function
def train(model, optimizer, epochs, device, train_ds, train_loss, train_accuracy, valid_ds, val_loss, val_accuracy):
    """
    Performs the entire training loop. Prints the loss and accuracy per step and epoch.

    Args:
        model (keras Model) -- image classifier
        optimizer (keras Optimizer) -- optimizer to use during training
```

```

epochs (int) -- number of epochs
train_ds (tf Dataset) -- the train set containing image-label pairs
train_loss (keras Loss) -- loss function for training
train_accuracy (keras Metric) -- accuracy metric for training
valid_ds (Tensor) -- the val set containing image-label pairs
val_loss (keras Loss) -- loss object for validation
val_accuracy (keras Metric) -- accuracy metric for validation
'''
step = 0
loss = 0.0
for epoch in range(epochs):
    for x, y in train_ds:
        # training step number increments at each iteration
        step += 1
        with tf.device(device_name=device):
            ### START CODE HERE ###
            # Run one training step by passing appropriate model parameters
            # required by the function and finally get the loss to report the results
            loss = train_one_step(model, optimizer, x, y, train_loss, train_accuracy)
            ### END CODE HERE ###
        # Use tf.print to report your results.
        # Print the training step number, loss and accuracy
        tf.print('Step', step,
                ': train loss', loss,
                '; train accuracy', train_accuracy.result())

    with tf.device(device_name=device):
        for x, y in valid_ds:
            # Call the model on the batches of inputs x and get the predictions
            y_pred = model(x)
            loss = val_loss(y, y_pred)
            val_accuracy.update_state(y, y_pred)

        # Print the validation loss and accuracy
        ### START CODE HERE ###
        tf.print('val loss', loss, '; val accuracy', val_accuracy.result())
        ### END CODE HERE ###

```

Run the `train` function to train your model! You should see the loss generally decreasing and the accuracy increasing.

Note: Please let the training finish before submitting and do not modify the next cell. It is required for grading. This will take around 5 minutes to run.

In [21]:

```

train(model, optimizer, EPOCHS, device, train_ds, train_loss, train_accuracy, valid_ds, val_loss, val_accuracy)

```

```

Step 1 : train loss 0.768829107 ; train accuracy 0.46875
Step 2 : train loss 0.510346174 ; train accuracy 0.5625
Step 3 : train loss 0.403072536 ; train accuracy 0.645833313
Step 4 : train loss 0.194760501 ; train accuracy 0.7265625
Step 5 : train loss 0.225985125 ; train accuracy 0.775
Step 6 : train loss 0.0859730393 ; train accuracy 0.8125
Step 7 : train loss 0.084870927 ; train accuracy 0.839285731
Step 8 : train loss 0.0551236 ; train accuracy 0.859375
Step 9 : train loss 0.0280453712 ; train accuracy 0.875
Step 10 : train loss 0.0530102961 ; train accuracy 0.8875
Step 11 : train loss 0.034141 ; train accuracy 0.897727251
Step 12 : train loss 0.0234619044 ; train accuracy 0.90625
Step 13 : train loss 0.0158608221 ; train accuracy 0.913461566
Step 14 : train loss 0.0189454015 ; train accuracy 0.919642866
Step 15 : train loss 0.0116201155 ; train accuracy 0.925
Step 16 : train loss 0.00849930942 ; train accuracy 0.9296875
Step 17 : train loss 0.00887948833 ; train accuracy 0.933823526
Step 18 : train loss 0.0939320475 ; train accuracy 0.935763896
Step 19 : train loss 0.00297227222 ; train accuracy 0.939144731
Step 20 : train loss 0.0162737016 ; train accuracy 0.942187488
Step 21 : train loss 0.00526584405 ; train accuracy 0.944940448
Step 22 : train loss 0.00630364101 ; train accuracy 0.947443187
Step 23 : train loss 0.00288483198 ; train accuracy 0.949728251
Step 24 : train loss 0.00214236067 ; train accuracy 0.951822937
Step 25 : train loss 0.00441805553 ; train accuracy 0.95375
Step 26 : train loss 0.0027152265 ; train accuracy 0.954987824

```

```

val loss 0.00354586402 ; val accuracy 1
Step 27 : train loss 0.00425177766 ; train accuracy 0.956674457
Step 28 : train loss 0.00392000657 ; train accuracy 0.958239257
Step 29 : train loss 0.00194190838 ; train accuracy 0.959695
Step 30 : train loss 0.00454552751 ; train accuracy 0.961052656
Step 31 : train loss 0.00356671098 ; train accuracy 0.962321818
Step 32 : train loss 0.00294305431 ; train accuracy 0.963510871
Step 33 : train loss 0.00181928789 ; train accuracy 0.964627147
Step 34 : train loss 0.00186028797 ; train accuracy 0.965677202
Step 35 : train loss 0.00171025819 ; train accuracy 0.966666639
Step 36 : train loss 0.00196678657 ; train accuracy 0.967600703
Step 37 : train loss 0.00205060793 ; train accuracy 0.968483806
Step 38 : train loss 0.00139280164 ; train accuracy 0.969320059
Step 39 : train loss 0.00477263331 ; train accuracy 0.970113099
Step 40 : train loss 0.00119135506 ; train accuracy 0.970866144
Step 41 : train loss 0.00254433742 ; train accuracy 0.971582174
Step 42 : train loss 0.00265221717 ; train accuracy 0.972263873
Step 43 : train loss 0.00174048753 ; train accuracy 0.972913623
Step 44 : train loss 0.00135703804 ; train accuracy 0.97353363
Step 45 : train loss 0.00157344353 ; train accuracy 0.974125862
Step 46 : train loss 0.050916627 ; train accuracy 0.974008203
Step 47 : train loss 0.00129105314 ; train accuracy 0.97456491
Step 48 : train loss 0.00235183211 ; train accuracy 0.975098312
Step 49 : train loss 0.00178711978 ; train accuracy 0.975609779
Step 50 : train loss 0.00240349839 ; train accuracy 0.976100624
Step 51 : train loss 0.0025420757 ; train accuracy 0.976572156
Step 52 : train loss 0.00244585727 ; train accuracy 0.976885617
val loss 0.00183823251 ; val accuracy 1

```

Evaluation

You can now see how your model performs on test images. First, let's load the test dataset and generate predictions:

In [22]:

```

test_imgs = []
test_labels = []

predictions = []
with tf.device(device_name=device):
    for images, labels in test_ds:
        preds = model(images)
        preds = preds.numpy()
        predictions.extend(preds)

        test_imgs.extend(images.numpy())
        test_labels.extend(labels.numpy())

```

Let's define a utility function for plotting an image and its prediction.

In [23]:

```

# Utilities for plotting

class_names = ['horse', 'human']

def plot_image(i, predictions_array, true_label, img):
    predictions_array, true_label, img = predictions_array[i], true_label[i], img[i]
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])

    img = np.squeeze(img)

    plt.imshow(img, cmap=plt.cm.binary)

    predicted_label = np.argmax(predictions_array)

    # green-colored annotations will mark correct predictions. red otherwise.
    if predicted_label == true_label:
        color = 'green'
    else:
        color = 'red'

```



```

color = red

# print the true label first
print(true_label)

# show the image and overlay the prediction
plt.xlabel("{} {:.20f}% ({})" .format(class_names[predicted_label],
                                     100*np.max(predictions_array),
                                     class_names[true_label]),
          color=color)

```

Plot the result of a single image

Choose an index and display the model's prediction for that image.

In [24]:

```

# Visualize the outputs

# you can modify the index value here from 0 to 255 to test different images
index = 8
plt.figure(figsize=(6,3))
plt.subplot(1,2,1)
plot_image(index, predictions, test_labels, test_imgs)
plt.show()

```

0



horse 100% (horse)

In []: