Mirrored Strategy: Basic

In this ungraded lab, you'll go through some of the basics of applying Mirrored Strategy.

Imports

In [1]:

```python
# Import TensorFlow and TensorFlow Datasets

import tensorflow_datasets as tfds
import tensorflow as tf
tfds.disable_progress_bar()

import os
```

Load the MNIST dataset and split it into training and test chunks.

In [2]:

```python
# Load the dataset we'll use for this lab
datasets, info = tfds.load(name='mnist', with_info=True, as_supervised=True, data_dir='./data')

mnist_train, mnist_test = datasets['train'], datasets['test']
```

Next, you define `strategy` using the `MirroredStrategy()` class. Print to see the number of devices available.

Note:

- If you are running this on Coursera, you'll see it gives a warning about no presence of GPU devices.
- If you are running this in Colab, make sure you have selected your `Runtime` to be `GPU` for it to detect it.
- In both these cases, you'll see there's only 1 device that is available.
- One device is sufficient for helping you understand these distribution strategies.

In [3]:

```python
# Define the strategy to use and print the number of devices found
strategy = tf.distribute.MirroredStrategy()
print('Number of devices: {}'.format(strategy.num_replicas_in_sync))
```

```
WARNING:tensorflow:There are non-GPU devices in `tf.distribute.Strategy`, not using nccl
allreduce.
```

```
WARNING:tensorflow:There are non-GPU devices in `tf.distribute.Strategy`, not using nccl
allreduce.
```

```
INFO:tensorflow:Using MirroredStrategy with devices
('/job:localhost/replica:0/task:0/device:CPU:0',)
```

```
INFO:tensorflow:Using MirroredStrategy with devices
('/job:localhost/replica:0/task:0/device:CPU:0',)
```

```
Number of devices: 1
```

Next, you create your training and test examples, define your batch size and also define `BATCH_SIZE_PER_REPLICA` which is the distribution you are making for each available device.

```python
# Get the number of examples in the train and test sets
num_train_examples = info.splits['train'].num_examples
num_test_examples = info.splits['test'].num_examples

BUFFER_SIZE = 10000

BATCH_SIZE_PER_REPLICA = 64
# Use for Mirrored Strategy
BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync
# Use for No Strategy
# BATCH_SIZE = BATCH_SIZE_PER_REPLICA * 1
```

A mapping function which normalizes your images:

```python
# Function for normalizing the image
def scale(image, label):
    image = tf.cast(image, tf.float32)
    image /= 255

    return image, label
```

Next, you create your training and evaluation datesets in the batch size you want by shuffling through your buffer size.

```python
# Set up the train and eval data set
train_dataset = mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
eval_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

For your model to follow the strategy, define your model within the strategy's scope.

- Run all the cells below and notice the results.
- Afterwards comment out `with strategy.scope():` and run everything again, without the strategy. Then you can compare the results. The important thing to notice and compare is the time taken for each epoch to complete. As mentioned in the lecture, doing a mirrored strategy on a single device (which are lab environment has) might take longer to train because of the overhead in implementing the strategy. With that, the advantages of using this strategy is more evident if you will use it on multiple devices.

```python
# Use for Mirrored Strategy -- comment out `with strategy.scope():` and deindent for no strategy
with strategy.scope():
    model = tf.keras.Sequential([
      tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),
      tf.keras.layers.MaxPooling2D(),
      tf.keras.layers.Flatten(),
      tf.keras.layers.Dense(64, activation='relu'),
      tf.keras.layers.Dense(10)
    ])
```

```python
model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              optimizer=tf.keras.optimizers.Adam(),
              metrics=['accuracy'])
```

```python
model.fit(train_dataset, epochs=12)
```

```
Epoch 1/12
WARNING:tensorflow:From /opt/conda/lib/python3.7/site-
packages/tensorflow/python/data/ops/multi_device_iterator_ops.py:601: get_next_as_optional (from t
ensorflow.python.data.ops.iterator_ops) is deprecated and will be removed in a future version.
Instructions for updating:
```

Use `tf.data.Iterator.get_next_as_optional()` instead.

```
938/938 [==============================] - 32s 34ms/step - loss: 0.2044 - accuracy: 0.9405
Epoch 2/12
938/938 [==============================] - 23s 25ms/step - loss: 0.0663 - accuracy: 0.9804
Epoch 3/12
938/938 [==============================] - 23s 25ms/step - loss: 0.0453 - accuracy: 0.9866
Epoch 4/12
938/938 [==============================] - 23s 25ms/step - loss: 0.0356 - accuracy: 0.9891
Epoch 5/12
938/938 [==============================] - 23s 25ms/step - loss: 0.0264 - accuracy: 0.9920
Epoch 6/12
938/938 [==============================] - 23s 25ms/step - loss: 0.0205 - accuracy: 0.9935
Epoch 7/12
938/938 [==============================] - 23s 25ms/step - loss: 0.0147 - accuracy: 0.9954
Epoch 8/12
938/938 [==============================] - 23s 24ms/step - loss: 0.0117 - accuracy: 0.9965
Epoch 9/12
938/938 [==============================] - 23s 24ms/step - loss: 0.0104 - accuracy: 0.9968
Epoch 10/12
938/938 [==============================] - 23s 24ms/step - loss: 0.0070 - accuracy: 0.9979
Epoch 11/12
938/938 [==============================] - 23s 24ms/step - loss: 0.0057 - accuracy: 0.9980
Epoch 12/12
938/938 [==============================] - 23s 24ms/step - loss: 0.0052 - accuracy: 0.9984
```

Out[9]:

<tensorflow.python.keras.callbacks.History at 0x7fbca07a2290>

In [ ]: