Open in Colab

One Device Strategy

In this ungraded lab, you'll learn how to set up a One Device Strategy. This is typically used to deliberately test your code on a single device. This can be used before switching to a different strategy that distributes across multiple devices. Please click on the **Open in Colab** badge above so you can download the datasets and use a GPU-enabled lab environment.

Imports

In [1]:

```
try:
  # %tensorflow_version only exists in Colab.
  %tensorflow_version 2.x
except Exception:
  pass

import tensorflow as tf
import tensorflow_hub as hub
import tensorflow_datasets as tfds

tfds.disable_progress_bar()
```

Define the Distribution Strategy

You can list available devices in your machine and specify a device type. This allows you to verify the device name to pass in `tf.distribute.OneDeviceStrategy()`.

In [2]:

```
# choose a device type such as CPU or GPU
devices = tf.config.list_physical_devices('GPU')
print(devices[0])

# You'll see that the name will look something like "/physical_device:GPU:0"
# Just take the GPU:0 part and use that as the name
gpu_name = "GPU:0"

# define the strategy and pass in the device name
one_strategy = tf.distribute.OneDeviceStrategy(device=gpu_name)
```

```
PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')
```

Parameters

We'll define a few global variables for setting up the model and dataset.

In [3]:

```
pixels = 224
MODULE_HANDLE = 'https://tfhub.dev/tensorflow/resnet_50/feature_vector/1'
IMAGE_SIZE = (pixels, pixels)
BATCH_SIZE = 32

print("Using {} with input size {}".format(MODULE_HANDLE, IMAGE_SIZE))
```

```
Using https://tfhub.dev/tensorflow/resnet_50/feature_vector/1 with input size (224, 224)
```

Download and Prepare the Dataset

Download and Prepare the Dataset

We will use the Cats vs Dogs dataset and we will fetch it via TFDS.

In [4]:

```
splits = ['train[:80%]', 'train[80%:90%]', 'train[90%:]']

(train_examples, validation_examples, test_examples), info = tfds.load('cats_vs_dogs', with_info=True, as_supervised=True, split=splits)

num_examples = info.splits['train'].num_examples
num_classes = info.features['label'].num_classes
```

Downloading and preparing dataset cats_vs_dogs/4.0.0 (download: 786.68 MiB, generated: Unknown size, total: 786.68 MiB) to /root/tensorflow_datasets/cats_vs_dogs/4.0.0...

```
WARNING:absl:1738 images were corrupted and were skipped
```

Shuffling and writing examples to /root/tensorflow_datasets/cats_vs_dogs/4.0.0.incompleteHDW9NW/cats_vs_dogs-train.tfrecord
Dataset cats_vs_dogs downloaded and prepared to /root/tensorflow_datasets/cats_vs_dogs/4.0.0. Subsequent calls will reuse this data.

In [5]:

```
# resize the image and normalize pixel values
def format_image(image, label):
    image = tf.image.resize(image, IMAGE_SIZE) / 255.0
    return  image, label
```

In [6]:

```
# prepare batches
train_batches = train_examples.shuffle(num_examples // 4).map(format_image).batch(BATCH_SIZE).prefetch(1)
validation_batches = validation_examples.map(format_image).batch(BATCH_SIZE).prefetch(1)
test_batches = test_examples.map(format_image).batch(1)
```

In [7]:

```
# check if the batches have the correct size and the images have the correct shape
for image_batch, label_batch in train_batches.take(1):
    pass

print(image_batch.shape)
```

```
(32, 224, 224, 3)
```

Define and Configure the Model

As with other strategies, setting up the model requires minimal code changes. Let's first define a utility function to build and compile the model.

In [8]:

```
# tells if we want to freeze the layer weights of our feature extractor during training
do_fine_tuning = False
```

In [9]:

```
def build_and_compile_model():
    print("Building model with", MODULE_HANDLE)

    # configures the feature extractor fetched from TF Hub
    feature_extractor = hub.KerasLayer(MODULE_HANDLE,
                                       input_shape=IMAGE_SIZE + (3,),
                                       trainable=do_fine_tuning)
```

```
    # define the model
    model = tf.keras.Sequential([
      feature_extractor,
      # append a dense with softmax for the number of classes
      tf.keras.layers.Dense(num_classes, activation='softmax')
    ])

    # display summary
    model.summary()

    # configure the optimizer, loss and metrics
    optimizer = tf.keras.optimizers.SGD(lr=0.002, momentum=0.9) if do_fine_tuning else 'adam'
    model.compile(optimizer=optimizer,
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    return model
```

You can now call the function under the strategy scope. This places variables and computations on the device you specified earlier.

In [10]:

```
# build and compile under the strategy scope
with one_strategy.scope():
    model = build_and_compile_model()
```

```
Building model with https://tfhub.dev/tensorflow/resnet_50/feature_vector/1
Model: "sequential"
_________________________________________________________________
Layer (type)                 Output Shape              Param #
=================================================================
keras_layer (KerasLayer)     (None, 2048)              23561152
_________________________________________________________________
dense (Dense)                (None, 2)                 4098
=================================================================
Total params: 23,565,250
Trainable params: 4,098
Non-trainable params: 23,561,152
_________________________________________________________________
```

`model.fit()` can be run as usual.

In [11]:

```
EPOCHS = 5
hist = model.fit(train_batches,
                 epochs=EPOCHS,
                 validation_data=validation_batches)
```

```
Epoch 1/5
582/582 [==============================] - 66s 113ms/step - loss: 0.0368 - accuracy: 0.9875 - val_
loss: 0.0254 - val_accuracy: 0.9923
Epoch 2/5
582/582 [==============================] - 65s 112ms/step - loss: 0.0207 - accuracy: 0.9936 - val_
loss: 0.0281 - val_accuracy: 0.9905
Epoch 3/5
582/582 [==============================] - 66s 113ms/step - loss: 0.0156 - accuracy: 0.9952 - val_
loss: 0.0281 - val_accuracy: 0.9923
Epoch 4/5
582/582 [==============================] - 66s 113ms/step - loss: 0.0115 - accuracy: 0.9963 - val_
loss: 0.0358 - val_accuracy: 0.9884
Epoch 5/5
582/582 [==============================] - 66s 113ms/step - loss: 0.0112 - accuracy: 0.9959 - val_
loss: 0.0299 - val_accuracy: 0.9936
```

Once everything is working correctly, you can switch to a different device or a different strategy that distributes to multiple devices.

In [ ]: